

The Mathematics of a Trading System

Expectancy, variance, the risk/reward tradeoff and the risk of ruin — measured, not assumed, on real OHLCV.

A walk through what a rules-based backtester actually proves once costs, position sizing and luck are taken seriously. Strategy: **ma_crossover**. Instruments: **SPY, QQQ, PETR4.SA, VALE3.SA, ITUB4.SA**.

Generated on 2026-06-20

1. Method — where the numbers come from

The win rate, expectancy and risk figures in this report are **not downloaded**. Yahoo Finance supplies only OHLCV; everything else is a *result* of simulating the strategy trade by trade. The engine is built to avoid the errors that flatter a backtest:


- **No lookahead.** A signal at the close of bar t is executed at the open of bar $t+1$ — never at the price that produced it.
- **Costs on every trade.** Spread, commission and slippage are charged round-trip; without them, paper expectancy lies.
- **Fixed-fractional risk.** Each trade risks a constant percentage of equity; a wider stop means a smaller position.
- **Conservative fills.** If a bar touches both stop and target, the stop is assumed to fill first.
- **Reproducible variance.** The Monte-Carlo bootstrap uses a fixed seed, so every run is identical.

Result in R (the R -multiple) is the unit that connects everything: a system is described by its win rate and its average R . A loss at the stop is $-1R$; the breakeven win rate, $1/(R+1)$, is the bar the real win rate must clear.


2. Cross-instrument scorecard

The same strategy, same parameters, run on every instrument. Where the expectancy is positive (after costs) the edge survives; where it is negative the geometry is just folklore on that market.

Instrument	Trades	Win%	Exp. R	Exp. \$	Profit factor	Max DD%	Edge?
SPY	39	30.8	-0.088	-4.61	0.87	5.6	no
QQQ	37	51.4	+0.422	+21.63	1.83	1.9	YES
PETR4.SA	42	35.7	-0.109	-5.63	0.84	3.7	no
VALE3.SA	43	41.9	+0.039	+1.70	1.06	3.6	YES
ITUB4.SA	40	47.5	+0.222	+11.09	1.40	4.7	YES



Per-trade expectancy in R, after costs. Green bars have a positive edge.



Actual win rate against the breakeven it must beat. A bar to the right of its grey twin is profitable.

But none of these point estimates is established. With ~40 trades each, the spread of expectancies is exactly what pure chance produces even if the true edge were identical (or zero) everywhere. The next section tests that honestly.

3. SPY — the full fingerprint


Metric	Value	Metric	Value
Trades	39	Profit factor	0.87
Win rate	30.77%	Payoff (R)	1.98
Avg win	101.45 (+2.06R)	Avg loss	-51.76 (-1.04R)
Expectancy \$	-4.61	Expectancy R	-0.088
Breakeven WR	33.59%	Actual WR	30.77%
Total return	-1.80%	CAGR	-0.12%
Max drawdown	5.58%	Sharpe / Sortino	-0.10 / -6.90

■ Small sample (39 trades < 100): expectancy is still noise-dominated. Treat with skepticism.

Equity curve — SPY




Underwater plot (drawdown) — SPY




3.1 Variance — same edge, different luck

Reshuffling the realized trades 5,000 times shows the range of outcomes the same system could have produced. The median is the honest expectation; the 5–95% band is the experience you must be able to sit through.




Monte-Carlo fan: median path and 5–95% band vs the single realized run.




Distribution of final equity across the simulations.

3.2 Risk of ruin & sample size



Probability of a deep drawdown by risk per trade — it climbs non-linearly as the risk dial turns up.




Running expectancy vs trade count: one trade means nothing; the average only settles with sample size.


4. QQQ — the full fingerprint

Metric	Value	Metric	Value
Trades	37	Profit factor	1.83
Win rate	51.35%	Payoff (R)	1.74
Avg win	92.63 (+1.81R)	Avg loss	-53.31 (-1.04R)
Expectancy \$	21.63	Expectancy R	+0.422
Breakeven WR	36.54%	Actual WR	51.35%
Total return	+8.00%	CAGR	+0.50%
Max drawdown	1.89%	Sharpe / Sortino	0.42 / 20.79

■ Small sample (37 trades < 100): expectancy is still noise-dominated. Treat with skepticism.




Realized equity curve, costs included.




Underwater plot: time spent below the previous equity peak.

4.1 Variance — same edge, different luck

Reshuffling the realized trades 5,000 times shows the range of outcomes the same system could have produced. The median is the honest expectation; the 5–95% band is the experience you must be able to sit through.




Monte-Carlo fan: median path and 5–95% band vs the single realized run.




Distribution of final equity across the simulations.

4.2 Risk of ruin & sample size



Probability of a deep drawdown by risk per trade — it climbs non-linearly as the risk dial turns up.



Running expectancy vs trade count: one trade means nothing; the average only settles with sample size.

5. PETR4.SA — the full fingerprint

Metric	Value	Metric	Value
Trades	42	Profit factor	0.84
Win rate	35.71%	Payoff (R)	1.52
Avg win	80.71 (+1.63R)	Avg loss	-53.60 (-1.08R)
Expectancy \$	-5.63	Expectancy R	-0.109
Breakeven WR	39.75%	Actual WR	35.71%
Total return	-2.37%	CAGR	-0.16%
Max drawdown	3.68%	Sharpe / Sortino	-0.14 / -4.86


■ Small sample (42 trades < 100): expectancy is still noise-dominated. Treat with skepticism.

Equity curve — PETR4.SA



Realized equity curve, costs included.


Underwater plot (drawdown) — PETR4.SA




Underwater plot: time spent below the previous equity peak.

5.1 Variance — same edge, different luck

Reshuffling the realized trades 5,000 times shows the range of outcomes the same system could have produced. The median is the honest expectation; the 5–95% band is the experience you must be able to sit through.




Monte-Carlo fan: median path and 5–95% band vs the single realized run.




Distribution of final equity across the simulations.

5.2 Risk of ruin & sample size



Probability of a deep drawdown by risk per trade — it climbs non-linearly as the risk dial turns up.




Running expectancy vs trade count: one trade means nothing; the average only settles with sample size.

6. VALE3.SA — the full fingerprint

Metric	Value	Metric	Value
Trades	43	Profit factor	1.06
Win rate	41.86%	Payoff (R)	1.48
Avg win	77.00 (+1.56R)	Avg loss	-52.51 (-1.06R)
Expectancy \$	1.70	Expectancy R	+0.039
Breakeven WR	40.37%	Actual WR	41.86%
Total return	+0.73%	CAGR	+0.05%
Max drawdown	3.62%	Sharpe / Sortino	0.04 / 2.40


■ Small sample (43 trades < 100): expectancy is still noise-dominated. Treat with skepticism.

Equity curve — VALE3.SA



Realized equity curve, costs included.


Underwater plot (drawdown) — VALE3.SA




Underwater plot: time spent below the previous equity peak.

6.1 Variance — same edge, different luck

Reshuffling the realized trades 5,000 times shows the range of outcomes the same system could have produced. The median is the honest expectation; the 5–95% band is the experience you must be able to sit through.





Monte-Carlo fan: median path and 5–95% band vs the single realized run.



Distribution of final equity across the simulations.

6.2 Risk of ruin & sample size



7. ITUB4.SA — the full fingerprint

Metric	Value	Metric	Value
Trades	40	Profit factor	1.40
Win rate	47.50%	Payoff (R)	1.54
Avg win	82.28 (+1.65R)	Avg loss	-53.33 (-1.07R)
Expectancy \$	11.09	Expectancy R	+0.222
Breakeven WR	39.32%	Actual WR	47.50%
Total return	+4.43%	CAGR	+0.28%
Max drawdown	4.70%	Sharpe / Sortino	0.25 / 10.54


■ Small sample (40 trades < 100): expectancy is still noise-dominated. Treat with skepticism.

Equity curve — ITUB4.SA



Realized equity curve, costs included.

Underwater plot (drawdown) — ITUB4.SA




Underwater plot: time spent below the previous equity peak.

7.1 Variance — same edge, different luck


Reshuffling the realized trades 5,000 times shows the range of outcomes the same system could have produced. The median is the honest expectation; the 5–95% band is the experience you must be able to sit through.

Monte-Carlo fan — variance changes everything (ITUB4.SA)





Monte-Carlo fan: median path and 5–95% band vs the single realized run.

Distribution of final equity across 5,000 simulations (ITUB4.SA)




Distribution of final equity across the simulations.

7.2 Risk of ruin & sample size



8. Is the edge real? Significance, pooling and cost

A point estimate without an interval is a guess with a confident voice. Bootstrapping the per-trade expectancy gives a 95% confidence interval; where it straddles zero, a small edge cannot be told apart from no edge at all.



Expectancy with its 95% bootstrap CI. A grey interval crosses zero — undecidable at this sample size.

Instrument	Trades	Expectancy R	95% CI (R)	P(>0)	Verdict
SPY	39	-0.088	[-0.531, +0.402]	34%	indistinguishable from zero
QQQ	37	+0.422	[-0.068, +0.926]	95%	indistinguishable from zero
PETR4.SA	42	-0.109	[-0.507, +0.314]	30%	indistinguishable from zero
VALE3.SA	43	+0.039	[-0.368, +0.468]	57%	indistinguishable from zero
ITUB4.SA	40	+0.222	[-0.196, +0.651]	85%	indistinguishable from zero
POOLED	201	+0.090	[-0.106, +0.295]	81%	indistinguishable from zero

Of the 5 instruments, **5** have a confidence interval that includes zero: individually, the data cannot confirm an edge in either direction. This is the honest reading the headline scorecard hides.

Two caveats make the picture even more sober. **QQQ sits exactly on the boundary** ($P(\text{edge} > 0) = 95\%$), so calling it positive would be a one-tailed knife-edge, not a finding. And **multiple comparisons** bite: testing five instruments, the chance that at least one clears a 95% one-sided bar by luck alone is $1 - 0.95^5 \approx 23\%$. Finding one borderline name among five is roughly what pure noise predicts — another reason not to single QQQ out.


8.1 Pooling for power, and an out-of-sample split

Because R normalises each trade by its risk, the 201 trades from all instruments can be pooled into one stream — finally crossing the ~100-trade threshold. The pooled expectancy is **+0.090R**. But these trades are *not* independent — US indices (~0.9 correlated) and the Brazilian names move together, so trades firing in the same period carry redundant information. An i.i.d. bootstrap would understate the uncertainty; a **calendar-quarter block bootstrap** keeps that correlation intact and gives the honest interval:

- i.i.d. CI (optimistic): [-0.106, +0.295]R
- block CI (58 quarters, honest): **[-0.143, +0.340]R** — still includes zero ($P(\text{edge} > 0) = 77\%$)

Splitting the pool chronologically into halves gives a genuine out-of-sample check:


- In-sample (first 100 trades): **-0.137R**
- Out-of-sample (last 101 trades): **+0.316R**



Pooled running expectancy. Even at ~200 trades it is still settling; the shaded zone marks the sub-100 region where the estimate is noise.

8.2 The edge is thin: cost sensitivity

Where expectancy is a fraction of an R, the cost assumption is a lever, not a footnote. Sweeping the slippage assumption shows how fast each instrument's edge crosses into the red.



Expectancy vs slippage. Lines that dip below the dashed zero line have lost their edge to costs — for the marginal names, that happens with a small change in assumptions.

9. The recovery math — why protecting capital wins

Losses and the gains needed to undo them are asymmetric. A 50% drawdown does not need 50% to recover — it needs 100%. This is why risk of ruin, not raw return, is the metric that keeps a trader in the game.

Drawdown suffered	Gain required to recover
10%	11%
20%	25%
30%	43%
40%	67%
50%	100%

Disclaimer

This is an educational and research project. Nothing here is investment advice. The example strategy is a didactic moving-average crossover, not a system with a guaranteed edge. The backtester *measures* an edge; it does not create one. Past performance does not guarantee future results.